

SentArg: A Hybrid Doc2Vec/DPH Model with Sentiment Analysis Refinement

Notebook for Touché: Argument Retrieval at CLEF 2020

Christian Staudte¹ and Lucas Lange²

¹ Leipzig University, Germany
cs47rica@studserv.uni-leipzig.de

² Leipzig University, Germany
l195wyqa@studserv.uni-leipzig.de

Abstract In this work we explore the yet untested inclusion of sentiment analysis in the argument ranking process. By comparing approaches, we show that favoring stronger sentiments can improve retrieval quality.

1 Introduction

In this notebook we propose our *SentArg* model for the first task of *Touché 2020: 1st Shared Task on Argument Retrieval*³ [6]. The task is performed on the *args.me corpus* [1]. The source code is available on our public GitHub repository⁴. *SentArg* is based on the *Dual Embedding Space Model* by Mitra et al. [11], although we replace BM25 with DPH [2] and adjust their scoring function. For further refinement we use sentiment analysis to classify the emotional level (positive, neutral, negative) of every argument utilizing the *Google Cloud Natural Language API*⁵. Analysing the author’s attitude, we thus re-rank arguments. Our main objective is to find out whether sentiments can increase the quality of an argument ranking.

The remainder of our notebook is organized as follows. Section 2 outlines all related works we used as theoretical and practical foundation. In Section 3 we describe our *SentArg* model configuration. Our evaluation as well as experiments are then presented in Section 4. Section 5 concludes our notebook.

2 Related Work

Words used in a mathematical model must be mapped onto mathematical objects. *One-hot* encodings are particularly well-known mappings for this: Each word is represented

Copyright © 2020 for this paper by its authors. Use permitted under Creative Commons License Attribution 4.0 International (CC BY 4.0). CLEF 2020, 22-25 September 2020, Thessaloniki, Greece.

³ <https://touche.webis.de>

⁴ <https://github.com/luckyos-code/ArgU>

⁵ <https://cloud.google.com/natural-language/>

by an element = 1 in vector $\vec{v} \in \{0,1\}^{|Vocab|}$, whereas all other elements are 0. However, these representations scale with vocabulary size and are therefore extremely sparse. Further, semantic relations between words are not modeled [7]. To overcome these limitations, Mikolov et al. [10] introduce two neural network models: *Continuous Bag-of-Words* (CBOW) and *Skip-Gram* (SG)⁶. Both architectures typically apply one input-, one hidden-, and one output-layer. For any given word w , its left- and right-sided context serves as input for a CBOW model, which tries to predict w . After training, weight matrix $W_{IN} \in \mathbf{R}^{M \times N}$, with N being the embedding's vector size, contains all available word embeddings. Words that are semantically similar also occur in similar contexts, hence cosine similarities of their embeddings are close to 1.

Regarding argument detection, similar contexts between queries and arguments may not be enough. To further restrict argument result sets, Mitra et al. [11] suggest a Dual Embedding Space Model. They calculate BM25 scores and word embedding similarities for any query Q regarding all arguments. Afterwards they form a weighted sum of both scores, which acts as argument relevance score. One special feature is the usage of CBOW's W_{OUT} matrix. By calculating the cosine similarity of two word embeddings, each from a different weight matrix, co-occurring words reach scores close to 1. In their final results, IN-OUT-similarities provide the best results.

Regarding typical information retrieval models, Potthast et al. [13] compare BM25, TFIDF, DirichletLM and DPH on more than 300.000 arguments (documents), containing 40 topics. Assessors evaluated different query results and analyzed various aspects of found arguments, such as relevance, as well as rhetorical, logical and dialectical quality. Their evaluation shows that DirichletLM and DPH are superior to BM25 and TFIDF. With these results in mind, we decide to replace BM25 in the Dual Embedding Space Model with DPH.

The goal of sentiment analysis is to identify the emotional character of a given text. This consideration is especially interesting in the case of arguments because research shows that arguments are fundamentally emotional [8,14]. The idea to include emotions in our ranking is based on the hypothesis that an emotional attachment to a topic can indicate involvement and thus stronger argumentation, which is backed by several studies [4,8,14,15]. We therefore expect a ranking in favor of emotional arguments to perform better than one favoring neutrality or one without sentiment analysis.

A survey conducted by Lawrence and Reed [9] shows that works on sentiment analysis in argument retrieval mainly focus on opinion mining, i.e. guessing the stance of an argument. Therefore they do not use an already annotated corpus, rather they instead aim to create one. Transitioning sentiment analysis from retrieving to ranking arguments is a new approach provided through our model.

3 Argument Retrieval Model

Our model features three steps: 1. Pre-processing; 2. Training of word embeddings and DPH to calculate argument and query similarities; 3. Sentiment analysis to re-sort arguments.

⁶ We primarily focus on CBOW in this notebook, so please refer to the referenced paper [10] for more information on SG.

Concerning Google’s sentiment analysis API, we first remove noise and formatting errors in as many arguments as possible. We manually define the following rule set:

- Remove URLs and square brackets with their content
- Remove some punctuation {~, #, §, &, @, =, *} }
- Reformat punctuation with correct spacing
- Replace identical letters in a row (>2) with one representative: *helloooooo* → *hello*

Short arguments often express approval or disapproval with previous arguments, hence we delete arguments containing less than 26 words. For our dual embedding’s CBOW model we replace/add the following rules:

- Numbers and URLs are replaced by tags <NUM> and <URL>
- Remove all punctuation
- Remove stop words

We keep upper- and lower-cases to reduce word ambiguities⁷ and train the CBOW model on all given arguments⁸. However, this gives rise to problems concerning rare query terms: "E-Cigarettes" occurs so rarely that no word embedding is trained. In this case we try different combinations of upper- and lower-case to find appropriate (and possibly multiple) word embeddings, which are all taken into account.

After preprocessing and training a CBOW model, we compute document vectors for each argument as proposed by [11]:

$$\bar{A} = \frac{1}{|A|} \cdot \sum_{a_j \in A} \frac{a_j}{||a_j||} \quad (1)$$

Each a_j represents the word embedding for word j . Afterwards we calculate similarities between a query and all arguments:

$$DESM(Q, A) = \frac{1}{|Q|} \cdot \sum_{q_i \in Q} \frac{q_i^T \bar{A}}{||q_i^T|| \cdot ||\bar{A}||} \quad (2)$$

In our experiments we test IN-IN and IN-OUT embeddings, but table 1 makes clear that IN-IN embeddings yield slightly better results for our approach. So both equations (1) and (2) only use this embedding type.

DPH is the second relevant component we model with the help of the Terrier framework [12]. Contrary to the original dual embedding architecture, we do not calculate a weighted sum of scores. Mitra et al. [11] state that word embeddings are helpful for finding appropriate contexts, whereas BM25 is better at finding concrete details given any query. We therefore run DESM and DPH in parallel, both select the top 1.000 arguments.⁹ Then, we form an intersection of both sets and sort every argument by its DPH score. Only context relevant scores thus have influence on the result set, while it is at the same time possible that a result set is empty.

⁷ For example "US" ≠ "us"

⁸ Configuration: vector size = 300; window size = 3; min word count = 5

⁹ A manually selected and adjusted parameter.

Sentiment analysis of the Google Cloud Natural Language API was created for a wide range of applications and is therefore most suitable for the diverse discussions in the data set. For every argument we send a request to the API, which in return provides us with a sentiment value. This value represents the direction of emotion on a scale from -1 to 1, with -1 and 1 being the strongest. Values between -0.1 and 0.1 represent neutral (N) sentiments, while values lower than -0.1 and greater than 0.1 express negative (-) and positive (+) sentiments, respectively.

Obtained sentiments are combined with DPH scores to compute the final ranking. To confirm our expectations we introduce two variants: (i) We encourage emotional arguments (values closer to -1 and 1); (ii) We encourage neutral arguments (values closer to 0). This is done using the following formula:

$$DPH_{Q,A}^* = DPH_{Q,A} \pm DPH_{Q,A} \cdot \frac{|S_A|}{2} \quad (3)$$

S_A refers to the sentiment value for argument A . Adding the weighted DPH refers to variant (i) and subtracting it to variant (ii). In the next section both approaches are compared to the retrieval result without applying sentiment analysis to test our hypothesis of favoring emotional arguments.

4 Experiments and Evaluation

When it comes to the embedding types IN-IN or IN-OUT, we manually evaluated the top 20 arguments for the first five queries respectively, see table 1. Each argument received a score between 0 (no relevance) and 3 (high relevance).¹⁰ On average IN-IN embeddings were slightly better, so we discarded the IN-OUT approach.

Since opinion mining states that sentiment analysis is capable of determining the stances of arguments [3,9], we tested if information is gained from sentiment analysis by comparing each argument’s stance with its sentiment value. In the dataset, each stance is stored as Boolean (pro and con), neutral stances are therefore not given. We enriched these stances by defining a sentiment value range that includes neutral (N), positive (+) and negative (-) arguments. Table 2 displays the distribution of sentiment values compared to annotated stances. As can be seen, most arguments (57%) are neutral, while 11% are positive and 32% are negative. However, positive and negative arguments do not match the dataset’s stance distribution of 53% pro and 47% con¹¹. Further, every sentiment class contains at least 42% candidates that are either pro or con. In spite of the opinion mining hypothesis, arguments for every sentiment class are nearly evenly distributed and no real correlation can be seen. In conclusion, by utilizing sentiment values we can gain new perspectives on the arguments.

Part of the task is the submission and evaluation of our model on the TIRA platform.¹² To identify the influence of our sentiment strategy on the retrieval, we evaluate

¹⁰ For more details, see our git repository’s experiments section <https://github.com/luckyos-code/ArgU/doc/experiments>.

¹¹ These numbers are not mentioned in Table 2, even though the distribution for neutral sentiments looks this way.

¹² <https://www.tira.io>; Group: *ir-lab-ul-tl-detroitnitz*

three different runs representing the different variants (see Section 3): R_0 No sentiment analysis, R_E Emotional is better [cp. (i)], and R_N Neutral is better [cp. (ii)]. The results are shown in Table 3. Compared to R_0 , our starting point before utilizing sentiments, R_N delivers the same scores for two metrics but offers no useful results in nDCG@5 and nDCG@10. Run R_E however gives the best results in all categories, confirming our expectations regarding emotional arguments performing better.

Table 1. Manually calculated Precision@20 and NDCG@20 for the first five queries. Q_1 : "Teachers Get Tenure", Q_2 : "Vaping E-Cigarettes Safe", Q_3 : "Insider Trading Allowed" Q_4 : "Corporal Punishment Used Schools", Q_5 : "Social Security Privatized". We set NDCG for Q_3 to 0 when calculating the average. For more details, visit our git repository’s evaluation section.

Emb.	Q_1		Q_2		Q_3		Q_4		Q_5		Avg.	
	Prec	NDCG	Prec	NDCG	Prec	NDCG	Prec	NDCG	Prec	NDCG	Prec	NDCG
IN-IN	0.45	0.94	0.35	0.82	0.0	0.36	0.9	0.97	0.9	0.92	0.52	0.80
IN-OUT	0.30	0.99	0.40	0.78	0.0	-	0.9	0.97	0.75	0.93	0.47	0.734

Table 2. Comparison of the sentiment (Sent.) values (N, +, -) of arguments and their stances (Pro, Con).

Sent.	Argument’s share in number (%)	Stance in % per share	
		Pro	Con
N	169,152 (57%)	53%	47%
+	33,406 (11%)	58%	42%
-	94,460 (32%)	50%	50%

Table 3. Results of the TIRA evaluation for runs R_0 , R_E and R_N .

	R_0	R_E	R_N
nDCG	0.223	0.232	0.223
nDCG@5	0.064	0.105	0.022
nDCG@10	0.073	0.105	0.058
QrelCoverage@10	0.32	0.46	0.32

5 Conclusion

By (1) combining word embeddings and DPH to retrieve relevant arguments and (2) improving our ranking by sorting arguments in regards to their sentiment value, we were able to show that favoring more emotional arguments can improve retrieval quality. Our empirical findings match the expectations set forth in our analysis of existing literature on the topic and enable us to confirm the potential of the SentArg model as a refinement strategy.

Future work may benefit from replacing the CBOW model with a FastText [5] architecture, which generates good results especially for noisy data. Further, static parameters and the influence equation for sentiments can be optimized.

References

1. Ajjour, Y., Wachsmuth, H., Kiesel, J., Potthast, M., Hagen, M., Stein, B.: Data Acquisition for Argument Search: The args.me corpus. In: Benz Müller, C., Stuckenschmidt, H. (eds.) 42nd German Conference on Artificial Intelligence (KI 2019). pp. 48–59. Springer (Sep 2019). https://doi.org/10.1007/978-3-030-30179-8_4
2. Amati, G.: Frequentist and bayesian approach to information retrieval. pp. 13–24 (01 1970). https://doi.org/10.1007/11735106_3
3. Bakshi, R.K., Kaur, N., Kaur, R., Kaur, G.: Opinion mining and sentiment analysis. In: 2016 3rd International Conference on Computing for Sustainable Global Development (INDIACom). pp. 452–455. IEEE (2016)
4. Blanchette, I., Caparos, S.: When emotions improve reasoning: The possible roles of relevance and utility. *Thinking & Reasoning* **19**(3-4), 399–413 (2013)
5. Bojanowski, P., Grave, E., Joulin, A., Mikolov, T.: Enriching word vectors with subword information. *Transactions of the Association for Computational Linguistics* **5**, 135–146 (2017)
6. Bondarenko, A., Hagen, M., Potthast, M., Wachsmuth, H., Beloucif, M., Biemann, C., Panchenko, A., Stein, B.: Touché: First Shared Task on Argument Retrieval. In: Castells, P., Ferro, N., Jose, J., Magalhães, J., Silva, M., Yilmaz, E. (eds.) *Advances in Information Retrieval. 42nd European Conference on IR Research (ECIR 2020)*. pp. 517–523. Lecture Notes in Computer Science, Springer, Berlin Heidelberg New York (Apr 2020). https://doi.org/10.1007/978-3-030-45442-5_67, https://link.springer.com/chapter/10.1007/978-3-030-45442-5_67
7. Braud, C., Denis, P.: Comparing word representations for implicit discourse relation classification. In: *Proceedings of the 2015 Conference on Empirical Methods in Natural Language Processing*. pp. 2201–2211. Association for Computational Linguistics, Lisbon, Portugal (Sep 2015). <https://doi.org/10.18653/v1/D15-1262>, <https://www.aclweb.org/anthology/D15-1262>
8. Gilbert, M.A.: *Coalescent argumentation*. Routledge (2013)
9. Lawrence, J., Reed, C.: Argument mining: A survey. *Computational Linguistics* **45**(4), 765–818 (2020)
10. Mikolov, T., Chen, K., Corrado, G., Dean, J.: Efficient estimation of word representations in vector space (2013)
11. Mitra, B., Nalisnick, E., Craswell, N., Caruana, R.: A dual embedding space model for document ranking (2016)
12. Ounis, I., Amati, G., Plachouras, V., He, B., Macdonald, C., Johnson, D.: Terrier information retrieval platform. In: Losada, D.E., Fernández-Luna, J.M. (eds.) *Advances in Information Retrieval*. pp. 517–519. Springer Berlin Heidelberg, Berlin, Heidelberg (2005)
13. Potthast, M., Gienapp, L., Euchner, F., Heilenkötter, N., Weidmann, N., Wachsmuth, H., Stein, B., Hagen, M.: Argument search: Assessing argument relevance. In: *Proceedings of the 42nd International ACM SIGIR Conference on Research and Development in Information Retrieval*. p. 1117–1120. SIGIR’19, Association for Computing Machinery, New York, NY, USA (2019). <https://doi.org/10.1145/3331184.3331327>, <https://doi.org/10.1145/3331184.3331327>
14. Villata, S., Cabrio, E., Jraidi, I., Benlamine, S., Chaouachi, M., Frasson, C., Gandon, F.: Emotions and personality traits in argumentation: an empirical evaluation 1. *Argument & Computation* **8**(1), 61–87 (2017)
15. Walton, D.: *The place of emotion in argument*. Penn State Press (2010)